

```
import kagglehub
aliiihussain_amazon_sales_dataset_path = kagglehub.dataset_download('aliiihussain/amazon-sales-dataset')

print('Data source import complete.')
```

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that gets preserved as output when you create a version using `kaggle competitions submit ...` or the UI
# You can also write temporary files to /kaggle/temp/, but they won't be saved outside of the current session
```

/kaggle/input/amazon-sales-dataset/amazon\_sales\_dataset.csv

## UNDERSTANDING THE DATA

```
import pandas as pd
import numpy as np

df=pd.read_csv("/kaggle/input/amazon-sales-dataset/amazon_sales_dataset.csv")
```

df

	order_id	order_date	product_id	product_category	price	discount_percent	quantity_sold	customer_region	payment_method
0	1	2022-04-13	2637	Books	128.75		10	4	North America
1	2	2023-03-12	2300	Fashion	302.60		20	5	Asia
2	3	2022-09-28	3670	Sports	495.80		20	2	Europe
3	4	2022-04-17	2522	Books	371.95		15	4	Middle East
4	5	2022-03-13	1717	Beauty	201.68		0	4	Middle East
...	...	...	...	...	...	...	...	...	...
49995	49996	2022-09-03	1433	Beauty	26.99		0	5	Middle East
49996	49997	2022-07-03	1428	Beauty	294.23		10	5	Asia
49997	49998	2023-02-17	4651	Electronics	352.11		30	4	Asia
49998	49999	2022-09-30	4371	Beauty	307.54		5	1	Middle East
49999	50000	2023-06-29	2944	Home & Kitchen	253.44		30	1	Europe

50000 rows × 13 columns

df.head()

	order_id	order_date	product_id	product_category	price	discount_percent	quantity_sold	customer_region	payment_method
0	1	2022-04-13	2637	Books	128.75		10	4	North America
1	2	2023-03-12	2300	Fashion	302.60		20	5	Asia
2	3	2022-09-28	3670	Sports	495.80		20	2	Europe
3	4	2022-04-17	2522	Books	371.95		15	4	Middle East
4	5	2022-03-13	1717	Beauty	201.68		0	4	Middle East

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   order_id          50000 non-null   int64  
 1   order_date        50000 non-null   object  
 2   product_id        50000 non-null   int64  
 3   product_category  50000 non-null   object  
 4   price             50000 non-null   float64 
 5   discount_percent 50000 non-null   int64  
 6   quantity_sold    50000 non-null   int64  
 7   customer_region   50000 non-null   object  
 8   payment_method    50000 non-null   object  
 9   rating            50000 non-null   float64 
 10  review_count     50000 non-null   int64  
 11  discounted_price 50000 non-null   float64 
 12  total_revenue    50000 non-null   float64 
dtypes: float64(4), int64(5), object(4)
memory usage: 5.0+ MB
```

```
import pandas as pd
df['order_date'] = pd.to_datetime(df['order_date'])
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   order_id          50000 non-null   int64  
 1   order_date        50000 non-null   object  
 2   product_id        50000 non-null   int64  
 3   product_category  50000 non-null   object  
 4   price             50000 non-null   float64 
 5   discount_percent 50000 non-null   int64  
 6   quantity_sold    50000 non-null   int64  
 7   customer_region   50000 non-null   object  
 8   payment_method    50000 non-null   object  
 9   rating            50000 non-null   float64 
 10  review_count     50000 non-null   int64  
 11  discounted_price 50000 non-null   float64 
 12  total_revenue    50000 non-null   float64 
dtypes: float64(4), int64(5), object(4)
memory usage: 5.0+ MB
```

```
df.isnull().sum()
```

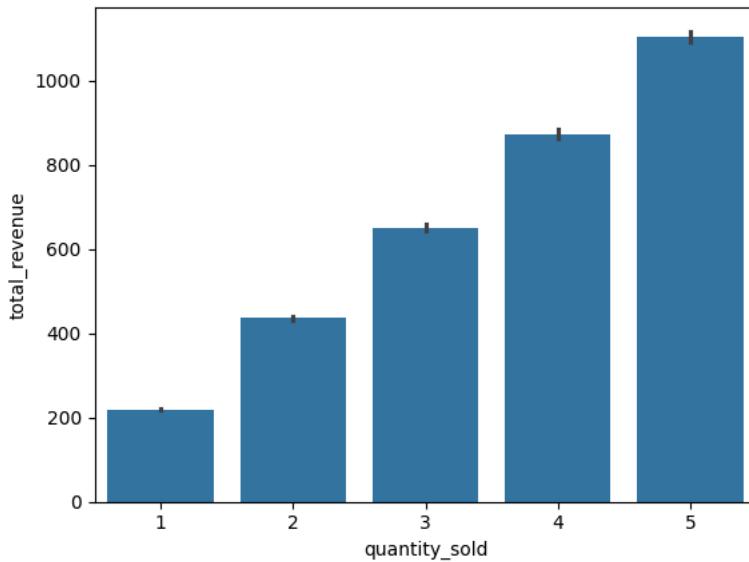
order_id	0
order_date	0
product_id	0
product_category	0
price	0
discount_percent	0
quantity_sold	0
customer_region	0
payment_method	0
rating	0
review_count	0
discounted_price	0
total_revenue	0
	dtype: int64

```
df['product_category'].value_counts()
```

product_category	
Beauty	8465
Fashion	8365
Books	8327
Electronics	8320
Sports	8265
Home & Kitchen	8258
Name: count, dtype: int64	

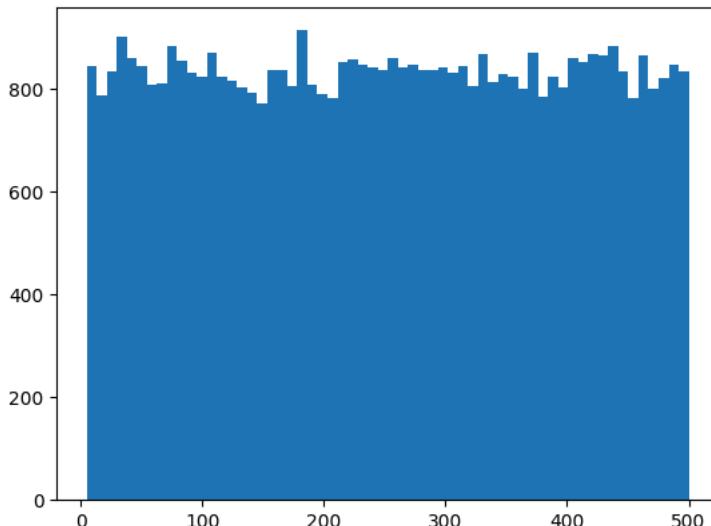
```
import seaborn as sns
import matplotlib.pyplot as plt
sns.barplot(x=df['quantity_sold'],y=df['total_revenue'] )
```

<Axes: xlabel='quantity\_sold', ylabel='total\_revenue'>



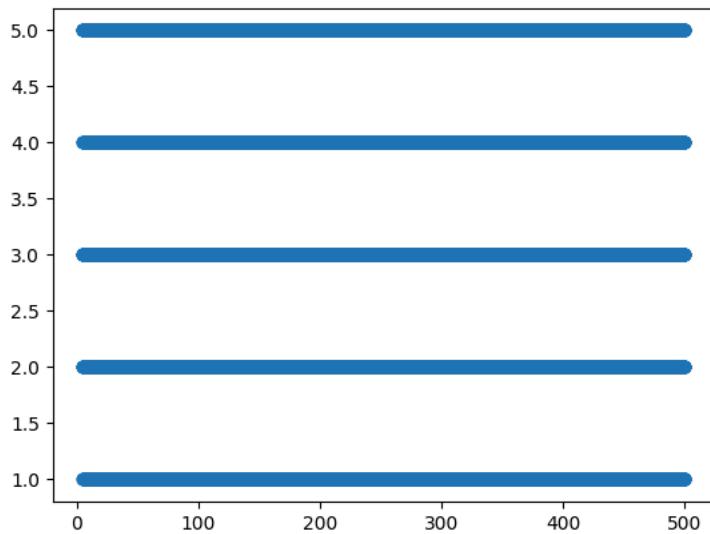
```
plt.hist(df['price'],bins=60)
```

```
(array([843., 786., 834., 902., 858., 844., 808., 810., 883., 854., 832.,
       824., 870., 824., 815., 802., 792., 772., 837., 837., 806., 913.,
       887., 789., 781., 851., 856., 847., 840., 836., 859., 841., 847.,
       835., 837., 840., 832., 844., 805., 866., 812., 828., 823., 799.,
       869., 785., 823., 802., 858., 851., 866., 864., 882., 833., 782.,
       865., 799., 821., 846., 833.]),
 array([
  5.01      , 13.25966667, 21.50933333, 29.759      ,
  38.00866667, 46.25833333, 54.508      , 62.75766667,
  71.00733333, 79.257      , 87.50666667, 95.75633333,
 104.006      , 112.25566667, 120.50533333, 128.755      ,
 137.00466667, 145.25433333, 153.504      , 161.75366667,
 170.00333333, 178.253      , 186.50266667, 194.75233333,
 203.002      , 211.25166667, 219.50133333, 227.751      ,
 236.00066667, 244.25033333, 252.5      , 260.74966667,
 268.99933333, 277.249      , 285.49866667, 293.74833333,
 301.998      , 310.24766667, 318.49733333, 326.747      ,
 334.99666667, 343.24633333, 351.496      , 359.74566667,
 367.99533333, 376.245      , 384.49466667, 392.74433333,
 400.994      , 409.24366667, 417.49333333, 425.743      ,
 433.99266667, 442.24233333, 450.492      , 458.74166667,
 466.99133333, 475.241      , 483.49066667, 491.74033333,
 499.99      ]),
<BarContainer object of 60 artists>)
```

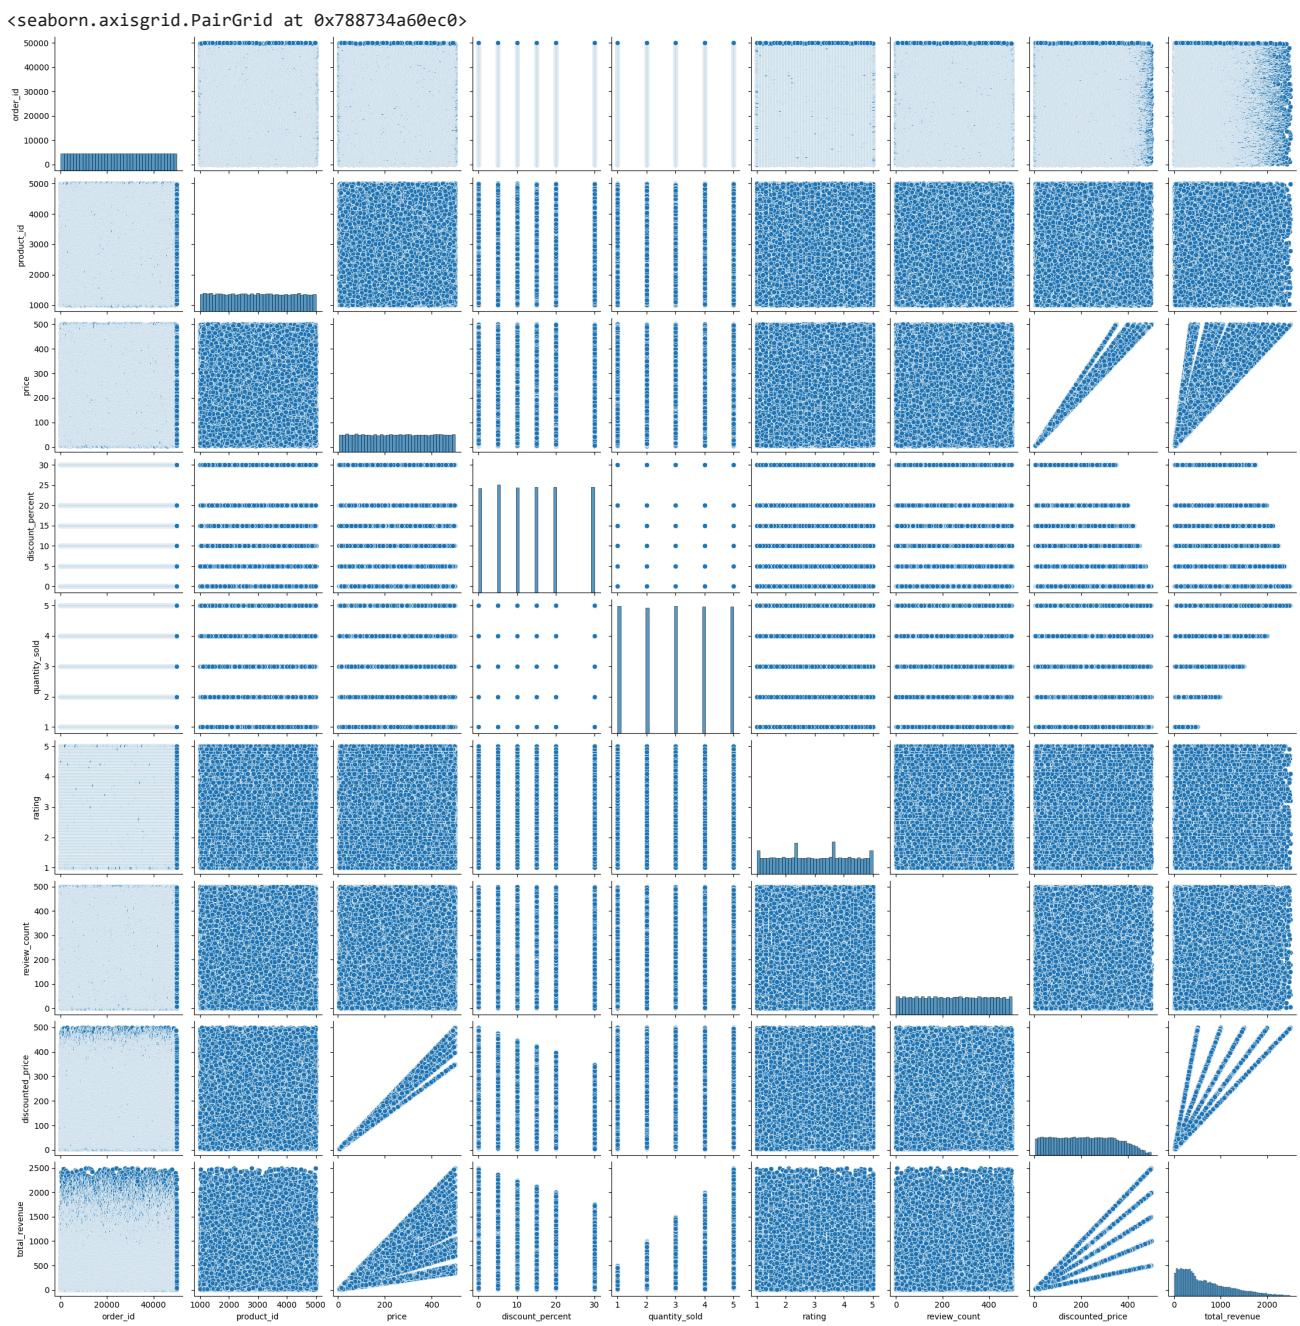


```
plt.scatter(df['price'],df['quantity_sold'])
```

```
<matplotlib.collections.PathCollection at 0x788734688110>
```



```
sns.pairplot(df)
```

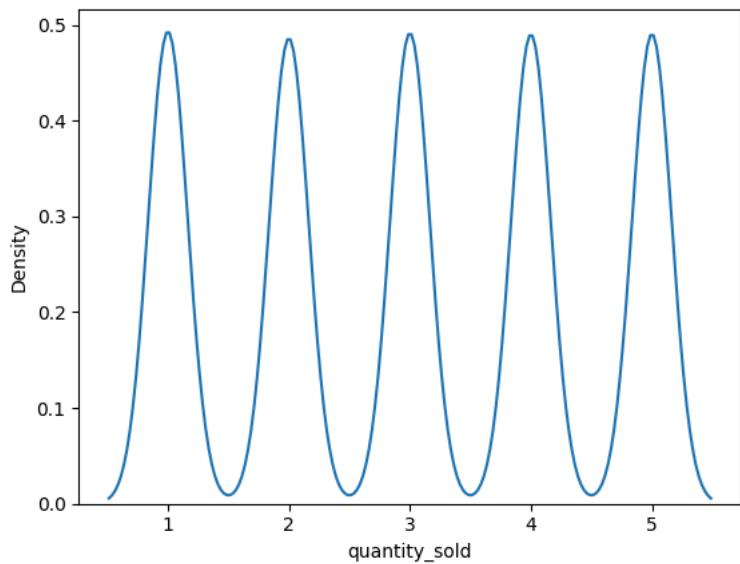


```
df.describe()
```

	order_id	product_id	price	discount_percent	quantity_sold	rating	review_count	discounted_price	to
<b>count</b>	500000.000000	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000
<b>mean</b>	25000.500000	2986.848740	252.507260	13.340700	2.999400	2.996316	249.329280	218.886566	5
<b>std</b>	14433.901067	1156.374535	143.025544	9.850694	1.415401	1.154295	144.251981	127.317681	
<b>min</b>	1.000000	1000.000000	5.010000	0.000000	1.000000	1.000000	0.000000	3.530000	
<b>25%</b>	12500.750000	1983.000000	127.840000	5.000000	2.000000	2.000000	125.000000	109.680000	
<b>50%</b>	25000.500000	2983.000000	252.970000	10.000000	3.000000	3.000000	250.000000	215.805000	
<b>75%</b>	37500.250000	3989.000000	376.335000	20.000000	4.000000	4.000000	374.000000	322.702500	
<b>max</b>	50000.000000	4999.000000	499.990000	30.000000	5.000000	5.000000	499.000000	499.910000	

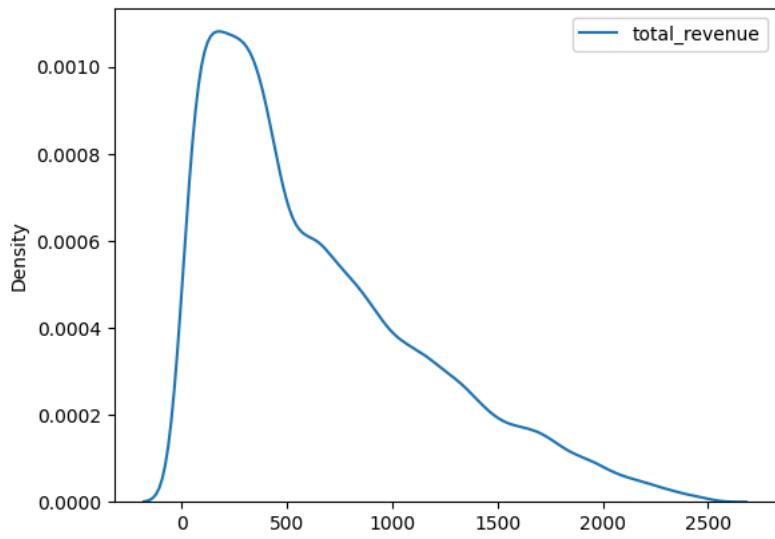
```
sns.kdeplot(df['quantity_sold'])
```

```
<Axes: xlabel='quantity_sold', ylabel='Density'>
```



```
sns.kdeplot([df['total_revenue']])
```

```
<Axes: ylabel='Density'>
```



## ▼ FEATURE TRANSFORMATION

### ▼ [A] FEATURE SCALING

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split( df[['quantity_sold', 'price']], df['rating'],test_size=0.2,random_state=10)
```

x\_train

	quantity_sold	price
20433	2	56.68
28866	2	44.59
42817	2	270.33
9543	2	416.77
20619	5	366.65
...	...	...
40059	3	66.31
28017	1	282.77
29199	4	157.15
40061	5	105.35
17673	1	249.09

40000 rows × 2 columns

## ▼ 1)STANDARDIZATION

```
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
scaler.fit(x_train)
x_train_scaled=scaler.transform(x_train)
x_test_scaled=scaler.transform(x_test)
```

df.describe()

	order_id	product_id	price	discount_percent	quantity_sold	rating	review_count	discounted_price	to
count	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000	5
mean	25000.500000	2986.848740	252.507260	13.340700	2.999400	2.996316	249.329280	218.886566	
std	14433.901067	1156.374535	143.025544	9.850694	1.415401	1.154295	144.251981	127.317681	
min	1.000000	1000.000000	5.010000	0.000000	1.000000	1.000000	0.000000	3.530000	
25%	12500.750000	1983.000000	127.840000	5.000000	2.000000	2.000000	125.000000	109.680000	
50%	25000.500000	2983.000000	252.970000	10.000000	3.000000	3.000000	250.000000	215.805000	
75%	37500.250000	3989.000000	376.335000	20.000000	4.000000	4.000000	374.000000	322.702500	
max	50000.000000	4999.000000	499.990000	30.000000	5.000000	5.000000	499.000000	499.910000	

## ▼ 2)Normalization

```
# min max scaling most useful
from sklearn.preprocessing import MinMaxScaler
scaler=MinMaxScaler()
scaler.fit(x_train)
x_train_scaler=scaler.transform(x_train)
x_test_scaler=scaler.transform(x_test)
```

df.describe()

	order_id	order_date	product_id	price	discount_percent	quantity_sold	rating	review_count	discounted_price
<b>count</b>	50000.000000		50000.000000	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000	50000.000000
<b>mean</b>	25000.500000	2022-12-31 10:47:16.800000256	2986.848740	252.507260	13.340700	2.999400	2.996316	249.329280	
<b>min</b>	1.000000	2022-01-01 00:00:00	1000.000000	5.010000	0.000000	1.000000	1.000000	0.000000	
<b>25%</b>	12500.750000	2022-07-02 00:00:00	1983.000000	127.840000	5.000000	2.000000	2.000000	125.000000	
<b>50%</b>	25000.500000	2023-01-02 00:00:00	2983.000000	252.970000	10.000000	3.000000	3.000000	250.000000	
<b>75%</b>	37500.250000	2023-07-02 00:00:00	3989.000000	376.335000	20.000000	4.000000	4.000000	374.000000	
<b>max</b>	50000.000000	2023-12-31 00:00:00	4999.000000	499.990000	30.000000	5.000000	5.000000	499.000000	

✓ [B] Encoding categorical `std: 1443.801867` `NaN` `1156.374535` `143.025544` `9.850694` `1.415401` `1.154295` `144.251981`

df									
The history saving thread hit an unexpected error (OperationalError('attempt to write a readonly database')).History will not be saved.									
	order_id	order_date	product_id	product_category	price	discount_percent	quantity_sold	customer_region	payment_method
0	1	2022-04-13	2637	Books	128.75		10	4	North America
1	2	2023-03-12	2300	Fashion	302.60		20	5	Asia
2	3	2022-09-28	3670	Sports	495.80		20	2	Europe
3	4	2022-04-17	2522	Books	371.95		15	4	Middle East
4	5	2022-03-13	1717	Beauty	201.68		0	4	Middle East
...	...	...	...	...	...	...	...	...	...
49995	49996	2022-09-03	1433	Beauty	26.99		0	5	Middle East
49996	49997	2022-07-03	1428	Beauty	294.23		10	5	Asia
49997	49998	2023-02-17	4651	Electronics	352.11		30	4	Asia
49998	49999	2022-09-30	4371	Beauty	307.54		5	1	Middle East
49999	50000	2023-06-29	2944	Home & Kitchen	253.44		30	1	Europe

50000 rows × 13 columns

df.isnull().sum()									
order_id	0								
order_date	0								
product_id	0								
product_category	0								
price	0								
discount_percent	0								
quantity_sold	0								
customer_region	0								
payment_method	0								
rating	0								
review_count	0								
discounted_price	0								
total_revenue	0								
dtype: int64									

df.info()									
<class 'pandas.core.frame.DataFrame'> RangeIndex: 50000 entries, 0 to 49999 Data columns (total 13 columns):									